

# 1. OpenAI — GPT-5.5

**Release:** 23 April 2026

**Knowledge Cutoff:** Nov 2025

**Context Window:** ~1.1 M tokens

**Pricing (API):**

Input: **\$5 / 1M tokens**

Output: **\$30 / 1M tokens**

**Availability:** Public (ChatGPT, API)

**Why / When to use**

Best-in-class **reasoning + agentic workflows**

Multi-step planning, coding agents, long-document analysis

High context + strong tool use

# 2. Google DeepMind — Gemini 3.1 Pro

**Release:** Feb 2026

**Knowledge Cutoff:** Mid-2025

**Context Window:** ~1 M tokens

**Pricing:** Usage-tiered (Vertex AI; lower than GPT-5.5 for long-context runs)

**Availability:** Public (Vertex AI, Gemini API)

**Why / When to use**

Strong **multimodal + long-context reasoning**

Competitive alternative to GPT-5.5 for enterprise analytics

Tight Google ecosystem  
integration

### 3. Anthropic — Claude Opus 4.7

**Release:** Mar 2026

**Knowledge Cutoff:** Oct 2025

**Context Window:** ~500K tokens

**Pricing:** Comparable to GPT-5 class (cheaper for tool-light workloads)

**Availability:** Public (API, Claude)

#### Why / When to use

- Safer, compliance-oriented reasoning
- Policy-sensitive enterprise use cases
- Strong long-form writing & analysis

### 4. Meta — Llama 3.1-405B

**Release:** July 2025

**Knowledge Cutoff:** Early 2025

**Context Window:** ~256K tokens

**Pricing:** Free (weights), inference cost dependent

**Availability:** Open-weight

#### Why / When to use

- On-prem or sovereign AI
- Custom fine-tuning
- Cost-optimized reasoning at scale

## 5. Mistral AI — Mistral Large 2

**Release:** Jan 2026

**Knowledge Cutoff:** 2025

**Context Window:** 128K tokens

**Pricing:** Lower than GPT-5.x tier

**Availability:** Public (API, Azure)

### Why / When to use

EU-friendly deployments

High-quality reasoning with lower latency

Strong coding + multilingual tasks

## 6. xAI — Grok-3

**Release:** Dec 2025

**Knowledge Cutoff:** Late 2025

**Context Window:** ~200K tokens

**Pricing:** Bundled / API limited

**Availability:** Partially public

### Why / When to use

Near-real-time reasoning with live data

Social and trend-aware analysis via X

Fast iteration workflows

## 7. Cohere — Command R+ (2026 refresh)

**Release:** Feb 2026

**Knowledge Cutoff:** 2025  
**Context Window:** 256K tokens  
**Pricing:** Enterprise-friendly  
**Availability:** Public (API)

**Why / When to use**

Retrieval-augmented generation (RAG)  
Search, enterprise QA  
Controlled, explainable reasoning

## 8. AI21 Labs — Jamba-Instruct 2

**Release:** Jan 2026  
**Knowledge Cutoff:** 2025  
**Context Window:** ~256K tokens  
**Pricing:** Competitive mid-tier  
**Availability:** Public (API)

**Why / When to use**

Hybrid SSM + Transformer efficiency  
Cost-effective long-context reasoning  
Document intelligence workloads

## 9. Amazon — Titan Text Premier (2026)

**Release:** Mar 2026  
**Knowledge Cutoff:** 2025  
**Context Window:** ~200K tokens  
**Pricing:** AWS-metered

**Availability:** Public (Bedrock)

**Why / When to use**

Native AWS integration  
Strict governance & IAM control  
Scalable enterprise reasoning

## 10. NVIDIA — Nemotron-4 340B

**Release:** Oct 2025

**Knowledge Cutoff:** 2025

**Context Window:** 128K tokens

**Pricing:** Open / enterprise support

**Availability:** Open-weight + NIMs

**Why / When to use**

On-device & data-center reasoning  
GPU-optimized efficiency  
Custom enterprise agents

